

HAN++ gap analysis for CareAI

April 10, 2026

Contents

HAN++ gap analysis for CareAI	2
What the current literature does support	2
Where HAN++ is most plausible	3
Main research gaps around HAN++	3
Lack of theory for superiority	3
Weak comparator landscape in many HAN style papers	3
Temporal modeling remains underdeveloped in classic HAN lines	4
Graph construction is under-justified	4
Interpretability claims are still weak	4
Safety claims outrun the evidence	5
External validation is still the main bottleneck	5
Joint diagnosis and confirmatory test recommendation is still thinly covered	5
Promising directions for future work	5
Direction 1	5
Direction 2	5
Direction 3	6
Direction 4	6
Direction 5	6
Direction 6	6
Direction 7	6
What evidence would be needed to make the HAN++ case convincing	7
Mathematical evidence	7
Methodological evidence	7
A realistic publication claim	7
Bottom line	8
References	9

HAN++ gap analysis for CareAI

CareAI sits in a crowded and fast-moving line of EHR prediction work. The existing project literature is enough to support a serious paper around HAN++, but not enough to support the strongest version of the claim that HAN++ is *the* best model class. The current evidence supports a narrower and more defensible position: a HAN style heterogeneous attention model is a credible choice for multi-entity clinical prediction, especially when interpretability, relation semantics, and sparse heterogeneous inputs matter [Hos18, Hos19, Qu22, Qu23]. The same literature also shows that stronger alternatives now exist along several axes, especially temporal modeling, structure learning, robustness to hub nodes, personalized knowledge integration, and multi-task training [Liu20, Cho19, Jia23b, Che24b, Cha24].

The most important gap is not just performance. It is evidence. The literature does not yet give a general mathematical proof that HAN style models are universally superior to traditional models or to newer graph and transformer variants. What it does provide is a set of local arguments and empirical patterns that can be turned into a strong paper if HAN++ is framed carefully and tested hard.

What the current literature does support

Question	What the literature supports	Strength of support	What is still missing
Is heterogeneous graph modeling reasonable for EHR diagnosis	Yes. Early and later work shows gains from modeling multi-type entities and relations rather than flattening records into simple vectors [Hos18, Hos19, Qu22, Sun20]	Moderate	External validation across hospitals and stronger non-graph baselines
Is hierarchical attention on heterogeneous medical graphs useful	Yes. Target-aware and semantic attention can improve diagnosis ranking and give more intuitive explanations [Hos19, Qu23, Yan23]	Moderate	Proof that the attention weights reflect causal or clinically faithful importance
Can uncertainty improve safety claims	Only partly. Uncertainty helps expose brittle cases, but population AUC alone is not enough and uncertainty by itself does not prove safety [Dus19]. Deferral methods can help, but require explicit evaluation of the human-AI workflow [Liu21]	Moderate	Real clinician workflow studies and patient-level decision analysis

Question	What the literature supports	Strength of support	What is still missing
Is small external validation enough for a strong clinical paper	Usually no. Calibration and net benefit often need much larger cohorts than many papers use [Ril21, Ram20, Cab21]	Strong	Adequate external cohorts with precise estimates
Can HAN++ be claimed as best overall	No clear support	Weak	Direct head-to-head evidence against modern graph, transformer, and strong non-graph models under one protocol

Where HAN++ is most plausible

HAN++ is most plausible when the problem has four features.

- The input is genuinely heterogeneous, with different node and edge types carrying different semantics [Hos18, Hos19].
- Meta-path structure captures clinically meaningful routes such as patient to symptom to disease or patient to test to disease [Qu22, Qu23].
- Attention is needed to expose which relation types and local neighborhoods drive a prediction [Hos19, Yan23].
- Data are sparse enough that relational inductive bias helps more than a flat model [Hos18, Cho16].

This is a good starting position for CareAI. It is not yet a winning position by itself.

Main research gaps around HAN++

Lack of theory for superiority

The current project literature contains almost no direct theory proving when a HAN style model should dominate simpler models, or when it should dominate newer graph architectures. Existing papers mostly argue from intuition and retrospective benchmark gains [Hos18, Hos19, Qu22]. That leaves three open questions.

- Which class of functions or clinical relations is HAN++ able to represent better than logistic regression, tree ensembles, sequence models, or non-attention graph models
- When does meta-path attention recover signal rather than amplify noise
- Under what graph conditions does heterogeneous attention remain stable under missingness, coding variation, and hospital shift

This is the single biggest gap if the goal is a mathematically strong paper.

Weak comparator landscape in many HAN style papers

Older heterogeneous diagnosis papers often beat weak baselines such as med2vec, shallow models, or limited RNNs [Hos18, Hos19]. That was enough at the time. It is not enough now. More recent work shows strong competing ideas.

Comparator class	Why it matters against HAN++	Representative paper
Graph structure learning	Learns hidden clinical structure instead of relying on fixed meta-path design	[Cho19]
Hub-aware heterogeneous graphs	Addresses a known failure mode of HAN on EHR graphs with high-degree nodes	[Liu20]
Personalized knowledge graph models	Inject outside biomedical knowledge and gain especially in low-data settings	[Jia23b]
Temporal heterogeneous graph transformers	Model both visit sequence and graph structure at once	[Che24b]
Multi-task heterogeneous models	Use related tasks to regularize representations and improve generalization	[Cha24]
Hypergraph and higher-order models	Capture code interactions beyond pairwise edges	[Xu23, Wan25]

A paper that does not beat or at least discuss these families will have trouble sustaining a strong novelty claim.

Temporal modeling remains underdeveloped in classic HAN lines

A major weakness of early heterogeneous attention diagnosis papers is that they focus on one admission or weak temporal summaries [Hos18, Hos19]. Newer methods show that explicit visit-to-visit state tracking matters [Che24b, Xu23b, Li25]. This creates a direct open question for CareAI.

Can HAN++ capture disease evolution well enough, or is it mainly a strong within-visit relational encoder that should be combined with a temporal module

This is both a gap and a promising future direction.

Graph construction is under-justified

Many medical graph papers assume that the graph schema, meta-paths, discretized labs, and relation definitions are correct, but offer limited sensitivity analysis [Hos18, Qu22, Liu20]. For CareAI, that leaves open questions about:

- which node types truly add value
- whether the chosen meta-paths are clinically grounded or just convenient
- how continuous labs should be encoded
- whether noisy hubs such as common symptoms or demographic nodes distort attention [Liu20]

A higher-class paper should treat graph construction as a scientific object, not just preprocessing.

Interpretability claims are still weak

Attention-based EHR papers often claim interpretability because they produce weights over nodes or semantic paths [Hos19, Qu23, Yan23]. That is useful, but still falls short of showing faithful explanation.

The literature still lacks strong tests of whether highlighted paths remain stable, clinically valid, and counterfactually meaningful. For CareAI, this means HAN++ explanations can be presented as heuristic evidence, not proof of mechanism, unless stronger faithfulness checks are added.

Safety claims outrun the evidence

The validation literature is clear that discrimination alone is not enough. Calibration, decision usefulness, and fit to the intended setting matter [Alb17, Vic06, Vic16]. The uncertainty literature is also clear that average test metrics can hide brittle patient-level decisions [Dus19]. Deferral can help, but only if the human-AI interaction is modeled and tested rather than assumed [Liu21].

This creates two open questions.

- Does HAN++ uncertainty identify genuinely dangerous cases, or merely low-confidence ones
- Does physician review after deferral improve the final decision enough to justify workflow cost

External validation is still the main bottleneck

For a clinical AI paper, a small external cohort may be enough to suggest transportability, but often not enough to estimate calibration slope and net benefit with useful precision [Ril21]. Methodology papers repeatedly warn that external validation should compare development and validation settings carefully and report uncertainty around performance estimates [Ram20, Cab21, Col15c].

If CareAI has only a small external set, that does not kill the paper. It does change the claim. The paper should then say the model has preliminary external evidence rather than strong generalizability proof.

Joint diagnosis and confirmatory test recommendation is still thinly covered

The project literature contains some work on lab recommendation and sequential diagnosis, but the evidence base is much thinner than for diagnosis prediction itself [Che18, Hua22, Raz23]. This is an opportunity. It is also a risk. The two-stage CareAI story may be novel, but the literature support is still less mature. That means the paper must be methodologically careful about what is claimed for the second stage.

Promising directions for future work

Direction 1

Build a conditional superiority claim rather than a universal best-model claim.

A stronger paper would argue that HAN++ is best suited for a defined regime:

- heterogeneous tabular EHR data
- explicit clinical entity relations
- moderate sample size
- need for path-level explanation
- multi-label disease prediction with sparse codes

That claim is much easier to defend than saying HAN++ is best in general.

Direction 2

Turn HAN++ into a temporal heterogeneous model.

The literature trend is moving from static heterogeneous graphs to models that combine temporal state tracking with structured relation learning [Che24b, Xu23b, Li25]. A strong future version of CareAI would keep HAN++ as the local relational engine and add a visit-level temporal layer on top.

Direction 3

Study meta-path necessity rather than assuming it.

A publishable contribution could come from showing which meta-paths matter, when they matter, and when learned structure can replace them. This would directly connect HAN++ to stronger baselines such as graph structure learning [Cho19] and hub-aware graph construction [Liu20].

Direction 4

Move from attention visualization to explanation stress tests.

Useful additions would include:

- stability of attention across seeds and folds
- clinician rating of highlighted nodes and paths
- counterfactual removal tests
- consistency between explanation and performance drop after feature masking

This would strengthen a part of the HAN++ story that current literature often leaves weak.

Direction 5

Make uncertainty decision-linked.

The strongest uncertainty papers do not stop at calibration plots. They connect uncertainty to patient-level instability and downstream action [Dus19, Liu21]. Future work should test whether HAN++ uncertainty:

- predicts error on shifted data
- predicts when physician review changes the final decision
- improves net benefit when tied to deferral thresholds

Direction 6

Push beyond pairwise graphs.

Hypergraph and higher-order interaction models are emerging because many clinical code sets act jointly rather than pairwise [Xu23, Wan25]. If HAN++ struggles against these models, that itself will reveal where heterogeneous attention is sufficient and where higher-order structure is needed.

Direction 7

Use multi-task learning if CareAI aims to become a broader platform.

Single-task diagnosis prediction may not be the best use of EHR structure. Recent work suggests that shared learning across mortality, readmission, drug recommendation, and diagnosis can improve robustness [Cha24]. Even if the first paper stays single-task, this is an important roadmap direction.

What evidence would be needed to make the HAN++ case convincing

Mathematical evidence

A strong paper would need at least one of these.

- A formal proposition showing when heterogeneous semantic attention is more expressive than a chosen traditional baseline under a clear graph data model
- A stability analysis showing bounded sensitivity to graph perturbation, missing edges, or hub-node noise
- An identifiability or consistency style argument for semantic attention weights under stated assumptions
- A computational argument showing why the chosen HAN++ factorization is a better bias-variance tradeoff than heavier alternatives for the target sample regime

Without this, the mathematical claim should stay modest.

Methodological evidence

The empirical package should be harder than what many prior HAN papers used.

Evidence needed	Why it matters	Literature basis
Strong baseline set with modern graph and non-graph models	Avoids winning only against dated comparators	[Cho19, Liu20, Jia23b, Che24b, Cha24]
Multi-seed results with variance reporting	Attention models can look stable at population level while differing at patient level	[Dus19]
Calibration, not just AUROC and F1	Clinical risk models need probability quality, not only ranking	[Alb17]
Decision curve analysis or net benefit	Shows clinical usefulness rather than abstract discrimination	[Vic06, Vic16]
External validation with honest precision limits	Small cohorts often mislead, especially for calibration	[Ril21, Ram20]
Transparent reporting of data, preprocessing, predictors, missingness, and model details	Exposes hidden bias and reproducibility problems	[Col15c, Wol19]
Ablations on graph schema and meta-paths	Tests whether HAN++ is strong because of architecture or because of hand-crafted graph design	[Hos18, Liu20]
Workflow study for deferral or physician review	Safety requires evidence about the final human-AI system	[Liu21]

A realistic publication claim

Based on the current literature, the most defensible high-class claim is not that HAN++ is universally the best. A stronger claim is this:

HAN++ offers a clinically meaningful inductive bias for heterogeneous EHR diagnosis prediction and can outperform traditional and some modern baselines when relation semantics, sparse multi-type inputs, and explanation needs are central, but this superiority must be shown under rigorous comparator, calibration, external validation, and workflow-aware evaluation.

That claim is evidence-aligned. It is still ambitious. It can support a good paper.

Bottom line

The literature gives a good reason to keep HAN++ in play. It does not yet give a good reason to assume HAN++ will win automatically. The open space for a strong paper is not merely to report better scores. It is to show *why* heterogeneous semantic attention helps, *when* it helps, and *how much* of that advantage survives stronger baselines, temporal modeling, distribution shift, and clinical decision constraints [Liu20, Cho19, Dus19, Ril21, Che24b, Cha24].

If CareAI can supply that evidence, HAN++ can anchor a serious paper. If not, HAN++ is better framed as a strong component inside a broader, more rigorously validated system rather than as the final answer.

References

- [Hos18] A. Hosseini, T. Chen, W. Wu, Y. Sun, and M. Sarrafzadeh, “HeteroMed: Heterogeneous Information Network for Medical Diagnosis,” *Proceedings of the 27th ACM International Conference on Information and Knowledge Management*, Apr. 2018, doi: [10.1145/3269206.3271805](https://doi.org/10.1145/3269206.3271805).
- [Hos19] A. Hosseini, T. Davis, and M. Sarrafzadeh, “Hierarchical Target-Attentive Diagnosis Prediction in Heterogeneous Information Networks,” *2019 International Conference on Data Mining Workshops (ICDMW)*, pp. 949–957, Nov. 2019, doi: [10.1109/ICDMW.2019.00138](https://doi.org/10.1109/ICDMW.2019.00138).
- [Qu22] Z. Qu, L. Cui, and Y. Xu, “Disease Risk Prediction via Heterogeneous Graph Attention Networks,” *2022 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*, pp. 3385–3390, Dec. 2022, doi: [10.1109/BIBM55620.2022.9995491](https://doi.org/10.1109/BIBM55620.2022.9995491).
- [Qu23] Z. Qu, L. Cui, and Y. Xu, “DP-MHAN: A Disease Prediction Method Based on Metapath Aggregated Heterogeneous Graph Attention Networks,” *International Conference on Database Systems for Advanced Applications*, pp. 276–285, 2023, doi: [10.1007/978-3-031-30678-5_21](https://doi.org/10.1007/978-3-031-30678-5_21).
- [Liu20] Z. Liu, X. Li, H. Peng, L. He, and P. S. Yu, “Heterogeneous Similarity Graph Neural Network on Electronic Health Records,” *2020 IEEE International Conference on Big Data (Big Data)*, pp. 1196–1205, Dec. 2020, doi: [10.1109/BigData50022.2020.9377795](https://doi.org/10.1109/BigData50022.2020.9377795).
- [Cho19] E. Choi *et al.*, “Learning the Graphical Structure of Electronic Health Records with Graph Convolutional Transformer,” *AAAI Conference on Artificial Intelligence*, pp. 606–613, Jun. 2019, doi: [10.1609/AAAI.V34I01.5400](https://doi.org/10.1609/AAAI.V34I01.5400).
- [Jia23b] P. Jiang, C. Xiao, A. Cross, and J. Sun, “GraphCare: Enhancing Healthcare Predictions with Personalized Knowledge Graphs,” *International Conference on Learning Representations*, May 2023.
- [Che24b] J. Chen, C. Yin, Y. Wang, and P. Zhang, “Predictive Modeling with Temporal Graphical Representation on Electronic Health Records,” *IJCAI : proceedings of the conference*, vol. 2024, pp. 5763–5771, May 2024, doi: [10.48550/arXiv.2405.03943](https://doi.org/10.48550/arXiv.2405.03943).
- [Cha24] T. H. Chan, G. Yin, K. Bae, and L. Yu, “Multi-task heterogeneous graph learning on electronic health records,” *Neural networks : the official journal of the International Neural Network Society*, vol. 180, pp. 106644, Aug. 2024, doi: [10.1016/j.neunet.2024.106644](https://doi.org/10.1016/j.neunet.2024.106644).
- [Sun20] Z. Sun, H. Yin, H. Chen, T. Chen, L. Cui, and F. Yang, “Disease Prediction via Graph Neural Networks,” *IEEE Journal of Biomedical and Health Informatics*, vol. 25, pp. 818–826, Jun. 2020, doi: [10.1109/JBHI.2020.3004143](https://doi.org/10.1109/JBHI.2020.3004143).
- [Yan23] Z. Yan *et al.*, “MG-HetGNN: Metapath-Guided Heterogeneous Graph Neural Network with Semantic Aggregation for Medical Diagnosis,” Aug. 28, 2023. doi: [10.1109/SWC57546.2023.10448852](https://doi.org/10.1109/SWC57546.2023.10448852).
- [Dus19] M. W. Dusenberry *et al.*, “Analyzing the role of model uncertainty for electronic health records,” *Proceedings of the ACM Conference on Health, Inference, and Learning*, Jun. 2019, doi: [10.1145/3368555.3384457](https://doi.org/10.1145/3368555.3384457).
- [Liu21] J. Liu, B. Gallego, and S. Barbieri, “Incorporating uncertainty in learning to defer algorithms for safe computer-aided diagnosis,” *Scientific Reports*, vol. 12, Aug. 2021, doi: [10.1038/s41598-022-05725-7](https://doi.org/10.1038/s41598-022-05725-7).
- [Ril21] R. Riley *et al.*, “Minimum sample size for external validation of a clinical prediction model with a binary outcome,” *Statistics in Medicine*, vol. 40, pp. 4230–4251, May 2021, doi: [10.1002/sim.9025](https://doi.org/10.1002/sim.9025).

- [Ram20] C. L. Ramspek, K. Jager, F. Dekker, C. Zoccali, and M. van Diepen, “External validation of prognostic models: what, why, how, when and where?” *Clinical Kidney Journal*, vol. 14, pp. 49–58, Nov. 2020, doi: [10.1093/ckj/sfaa188](https://doi.org/10.1093/ckj/sfaa188).
- [Cab21] F. Cabitza *et al.*, “The importance of being external. methodological insights for the external validation of machine learning models in medicine,” *Computer methods and programs in biomedicine*, vol. 208, pp. 106288, Jul. 2021, doi: [10.1016/j.cmpb.2021.106288](https://doi.org/10.1016/j.cmpb.2021.106288).
- [Cho16] E. Choi, M. T. Bahadori, L. Song, W. Stewart, and J. Sun, “GRAM: Graph-based Attention Model for Healthcare Representation Learning,” *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, Nov. 2016, doi: [10.1145/3097983.3098126](https://doi.org/10.1145/3097983.3098126).
- [Xu23] R. Xu, M. K. Ali, J. Ho, and C. Yang, “Hypergraph Transformers for EHR-based Clinical Predictions.” *AMIA Joint Summits on Translational Science proceedings. AMIA Joint Summits on Translational Science*, vol. 2023, pp. 582–591, 2023.
- [Wan25] H. Wang and Y. Yuan, “Structure-aware Hypergraph Transformer for Diagnosis Prediction in Electronic Health Records,” *2025 IEEE International Conference on Systems, Man, and Cybernetics (SMC)*, pp. 1703–1708, Aug. 2025, doi: [10.1109/SMC58881.2025.11343341](https://doi.org/10.1109/SMC58881.2025.11343341).
- [Xu23b] Y. Xu *et al.*, “Time-Aware Context-Gated Graph Attention Network for Clinical Risk Prediction,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 35, pp. 7557–7568, Jul. 2023, doi: [10.1109/TKDE.2022.3181780](https://doi.org/10.1109/TKDE.2022.3181780).
- [Li25] D. Li, Z. Yao, M. Liang, and M. Liu, “DeepJ: Graph Convolutional Transformers with Differentiable Pooling for Patient Trajectory Modeling,” *AMIA ... Annual Symposium proceedings. AMIA Symposium*, vol. 2024, pp. 734–743, Jun. 2025, doi: [10.48550/arXiv.2506.15809](https://doi.org/10.48550/arXiv.2506.15809).
- [Alb17] A. Alba *et al.*, “Discrimination and Calibration of Clinical Prediction Models: Users’ Guides to the Medical Literature,” *JAMA*, vol. 318, pp. 1377–1384, Oct. 2017, doi: [10.1001/jama.2017.12126](https://doi.org/10.1001/jama.2017.12126).
- [Vic06] A. Vickers and E. Elkin, “Decision Curve Analysis: A Novel Method for Evaluating Prediction Models,” *Medical Decision Making*, vol. 26, pp. 565–574, Nov. 2006, doi: [10.1177/0272989X06295361](https://doi.org/10.1177/0272989X06295361).
- [Vic16] A. Vickers, B. V. calster, and E. Steyerberg, “Net benefit approaches to the evaluation of prediction models, molecular markers, and diagnostic tests,” *The BMJ*, vol. 352, Jan. 2016, doi: [10.1136/bmj.i6](https://doi.org/10.1136/bmj.i6).
- [Col15c] G. Collins, J. B. Reitsma, D. G. Altman, and K. Moons, “Transparent reporting of a multivariable prediction model for individual prognosis or diagnosis (TRIPOD): the TRIPOD Statement,” *BMC Medicine*, vol. 13, Jan. 2015, doi: [10.1186/s12916-014-0241-z](https://doi.org/10.1186/s12916-014-0241-z).
- [Che18] L.-F. Cheng, N. Prasad, and B. Engelhardt, “An Optimal Policy for Patient Laboratory Tests in Intensive Care Units,” *Pacific Symposium on Biocomputing. Pacific Symposium on Biocomputing*, vol. 24, pp. 320–331, Aug. 2018, doi: [10.1142/9789813279827_0029](https://doi.org/10.1142/9789813279827_0029).
- [Hua22] T. Huang, L. Li, E. Bernstam, and X. Jiang, “Confidence-based laboratory test reduction recommendation algorithm,” *BMC Medical Informatics and Decision Making*, vol. 23, Jun. 2022, doi: [10.1186/s12911-023-02187-3](https://doi.org/10.1186/s12911-023-02187-3).
- [Raz23] S. Raza and C. Ding, “Improving Clinical Decision Making With a Two-Stage Recommender System,” *IEEE/ACM Transactions on Computational Biology and Bioinformatics*, vol. 21, pp. 1180–1190, Sep. 2023, doi: [10.1109/TCBB.2023.3318209](https://doi.org/10.1109/TCBB.2023.3318209).

[Wol19] R. Wolff *et al.*, “PROBAST: A Tool to Assess the Risk of Bias and Applicability of Prediction Model Studies,” *Annals of Internal Medicine*, vol. 170, pp. 51–58, Jan. 2019, doi: [10.7326/M18-1376](https://doi.org/10.7326/M18-1376).